

Anthony Reyna

anthonyReynaX@gmail.com | linkedin.com/in/noStackEngineer | [Personal Website: anthony-reyna.com](https://anthony-reyna.com)

EDUCATION

University of California, Santa Cruz

Bachelor of Science in Computer Science

Santa Cruz, CA

Sep. 2023 – June 2026

EXPERIENCE

Software Engineering Intern (AI Systems)

June 2026 – Aug. 2026

EUV Tech

Martinez, CA

- Diagnosed a dashboard undercounting events by two orders of magnitude and replaced its client-side aggregation with a REST endpoint that fans out to 20+ tool databases and buckets events in SQL; fixed a request-parameter mismatch that had silently disabled its filters.
- Cut that timeline's page loads from ~40 s to under 1 s by directing Claude Code through profiling and the fix: keyset pagination across 20+ databases, missing indexes, and an hourly rollup refreshed by a bounded worker pool that flags any database that fails to refresh instead of silently undercounting.
- Moved an LLM assistant's engineering knowledge base from its always-loaded prompt (a context window at capacity) into a Neo4j knowledge graph of the equipment fleet's physical connections and layouts, which scales as new equipment comes online; added graph-traversal tools and reworked the prompt stack so the assistant uses them when a question needs the graph.
- Traced a production incident to its root cause across three layers of the stack – browser, nginx routing, and a hand-maintained service config – after a misrouted request returned a 200 status with the wrong content and no errors anywhere in the logs.
- Built an automated evaluation harness for a production LLM assistant, with tiered cost/token accounting, rubric-based scoring, and content-addressed state fingerprinting that make performance comparisons across prompt and model changes reproducible.
- Worked as a technical liaison between engineering and field service teams, translating operational needs into product features; consulted by project leads on infrastructure and evaluation practices.

Research Intern

June 2024 – Aug. 2024

Carnegie Mellon University

Pittsburgh, PA

- Evaluated LLM personal assistants for robustness and security risks using the Enron Email Dataset (500,000+ messages), and proposed mitigations for adversarial attacks on agent pipelines.

Undergraduate Researcher

Jan. 2024 – May 2024

University of California, Santa Cruz

Santa Cruz, CA

- Developed a benchmark suite spanning 4 multi-agent coordination games (Hanabi, Overcooked-AI, Collab Capture, Collab Escape) evaluating LLM agentic coordination; co-authored *LLM-Coordination*, accepted to NAACL 2025 Findings.

PROJECTS

Farming Prototype | *Unreal Engine 5.7, C++, Blueprints, PaperZD*

Sep. 2026 – Present

- Building a 2.5D farming game in Unreal Engine 5 with gameplay systems written in C++ and exposed to Blueprints: a tile-based farm grid with a per-tile state machine (till, plant, water, harvest) and a day cycle that drives crop growth.
- Designed a Blueprint-implementable interaction interface with priority-based target selection over collision overlap queries, letting the player interact with any world object through a single interface, so new interactables (including Blueprint-only ones) need no changes to the player code.
- Wrote editor tooling that fits the grid to level geometry using per-tile line traces, saving the results with the level so no traces run at game start, plus a toggleable grid preview and in-editor buttons for testing game logic.
- Batched tile and crop rendering with instanced static meshes and grouped sprites, reusing one instanced component per mesh and material pair; drove directional sprite animation from player input with Enhanced Input.

Multithreaded HTTP Server | *C, POSIX Threads, Sockets* – Principles of Computer Systems Design Winter 2026

- Built an HTTP/1.1 file server in C on raw sockets that reassembles requests from partial reads, rejects malformed ones, and streams Content-Length bodies to disk; PUTs are received outside any lock and committed atomically by rename, so slow clients never block other writers.

- Made it concurrent with a dispatcher thread feeding a configurable worker pool through a bounded blocking queue (mutex + condition variables); wrote reader-writer locks with reader, writer and n-way priority, and a hashed per-file lock map so requests to different files never contend.

TECHNICAL SKILLS

Languages: Python, C/C++, TypeScript, SQL (PostgreSQL), Java

AI-Assisted Development: Claude Code (directing, reviewing and debugging agent-written code), multi-agent team workflows (file-based task locks, shared work logs)

AI/ML: LLMs, multimodal LLMs, VLA models, tokenizers, fine-tuning & inference on GPU clusters, LLM evaluation & benchmarking, PyTorch, Hugging Face, AWS Bedrock

Game Development: Unreal Engine 5 (C++, Blueprints, Enhanced Input, Paper2D/PaperZD), Git LFS

Systems & Tools: Concurrency (POSIX threads), HTTP & Sockets, REST APIs, Data Structures & Algorithms, OOP, Neo4j, Git, Linux